

Mid-Term Practice Portion

Problem-1: User interaction with an IR system

Pick a site with a good number of text, html, or pdf documents. Crawl it using wget, collecting 100-1000 documents. You may have to do some trial-and-error to get to this number. Make sure you don't keep your crawler running for too long. You will be penalized for having fewer than 100 or more than 1000 indexed documents. Report your wget command. [5 points]

```
# This is my wget command that I used

# wget -r paulgraham.com
```

Index this collection using Pyterrier. Make sure you pick appropriate meta tags that may be useful for retrieval later. Report the index statistics. [5 points]

```
import pyterrier as pt
if not pt.started():
    pt.init()
```

C:\Users\Fortu\AppData\Local\Temp\ipykernel_9280\1742394333.py:2: DeprecationWarning: Call to deprecated function (or staticmethod) started. (use pt.java.started() instead) -- Deprecated since version 0.11.0.

```
if not pt.started():
```

```
files = pt.io.find_files('c:/Users/Fortu/Downloads/Aut2024/Info 376/paulgraham.com')
indexer = pt.FilesIndexer('c:/Users/Fortu/Downloads/Aut2024/Info 376/midterm_practice/webcrawler_index',
                           verbose=True, blocks=False,
                           meta={"docno":20,"filename":1024,"title":1024},meta_tags={"title":"title"},
                           overwrite=True)
indexref = indexer.index(files)
index = pt.IndexFactory.of(indexref)
print(index.getCollectionStatistics().toString())
```

```
16:24:03.828 [main] ERROR org.terrier.structures.indexing.Indexer -- Could not rename index
java.io.IOException: Rename of index structure file 'c:/Users/Fortu/Downloads/Aut2024/Info
376/midterm_practice/webcrawler_index/data_1.direct.bf' (exists) to
'c:/Users/Fortu/Downloads/Aut2024/Info 376/midterm_practice/webcrawler_index/data.direct.bf'
(exists) failed - likely that source file is still open. Possible indexing bug?
    at org.terrier.structures.IndexUtil.renameIndex(IndexUtil.java:379)
    at org.terrier.structures.indexing.Indexer.index(Indexer.java:388)
```

Number of documents: 384

Number of terms: 12539

Number of postings: 160977

Number of fields: 0

Number of tokens: 412945

Field names: []

Positions: false

Create a client-side interface in PHP/HTML by making appropriate modifications to what was provided in the class. Make sure that in addition to the search box and button, it has some description of what the underlying collection is and offer some suggestions for what one could search for. [5 points]

You can access the search interface using the following link: [Search Interface](#)

(<http://is-searchrec.ischool.uw.edu/swas/search.php>)

```
import pandas as pd
queries = pd.DataFrame([["q1", "addiction"], ["q2", "life"], ["q3", "age"], ["q4", "USA"]], columns=
queries
```

| | qid | query |
|---|-----|-----------|
| 0 | q1 | addiction |
| 1 | q2 | life |
| 2 | q3 | age |
| 3 | q4 | USA |

```
index = pt.IndexFactory.of("c:/Users/Fortu/Downloads/Aut2024/Info 376/midterm_practice/webcrawler
BM = pt.BatchRetrieve(index, wmodel="BM25")
BM.transform(queries)
```

C:\Users\Fortu\AppData\Local\Temp\ipykernel_9280\1970710156.py:2: DeprecationWarning: Call to deprecated class BatchRetrieve. (use pt.terrier.Retriever() instead) -- Deprecated since version 0.11.0.

```
BM = pt.BatchRetrieve(index, wmodel="BM25")
```

| | qid | docid | docno | rank | score | query |
|-----|-----|-------|-------|------|----------|-----------|
| 0 | q1 | 15 | d16 | 0 | 9.957912 | addiction |
| 1 | q1 | 87 | d88 | 1 | 8.445748 | addiction |
| 2 | q1 | 176 | d177 | 2 | 8.236193 | addiction |
| 3 | q1 | 352 | d353 | 3 | 8.205934 | addiction |
| 4 | q1 | 133 | d134 | 4 | 6.743001 | addiction |
| ... | ... | ... | ... | ... | ... | ... |
| 180 | q3 | 187 | d188 | 59 | 1.338348 | age |
| 181 | q3 | 316 | d317 | 60 | 1.030054 | age |

| | qid | docid | docno | rank | score | query |
|-----|-----|-------|-------|------|-----------|-------|
| 182 | q4 | 347 | d348 | 0 | 10.398220 | USA |
| 183 | q4 | 37 | d38 | 1 | 7.369159 | USA |
| 184 | q4 | 323 | d324 | 2 | 5.813612 | USA |

185 rows × 6 columns

Kill any existing server you have running in the background. Now edit and start the server script and have it listen to your assigned port. When one enters a query on your search page, it should display ranked results from your index. Submit the URL for this search page. [10 points]

```
# my assigned port (10014) was taken so I used 10001
```

Problem-2: Indexing to evaluation

Use the [LA Times data](#) in TREC format available from here to index using Pyterrier. Make sure you index with tags 'DOCNO' and 'HEADLINE'. Report index statistics. [10 points]

```
files = pt.io.find_files('c:/Users/Fortu/Downloads/Aut2024/Info 376/latimes')
indexer = pt.TRECCollectionIndexer('c:/Users/Fortu/Downloads/Aut2024/Info 376/midterm_practice/la',
                                   verbose=True, blocks=False,
                                   meta={"docno":20,"filename":1024,"headline":1024},meta_tags={"headline":"headline"},
                                   overwrite=True)
indexref = indexer.index(files)
index = pt.IndexFactory.of(indexref)
print(index.getCollectionStatistics().toString())
```

2files [01:46, 53.07s/files]

Number of documents: 131896

Number of terms: 188927

Number of postings: 24807798

Number of fields: 0

Number of tokens: 40507829

Field names: []

Positions: false

Show how you can retrieve the docno for a document in that collection by providing a docid (e.g., 0, 1, 2) with a couple of examples. Report your Pyterrier commands and outcomes. [5 points]

```
meta = index.getMetaIndex()
docid = 0
docno = meta.getItem("docno", docid)
print(f"Docid: {docid}, Docno: {docno}")

docid = 1
```

```
docno = meta.getItem("docno", docid)
print(f"Docid: {docid}, Docno: {docno}")
```

Docid: 0, Docno: LA010189-0001

Docid: 1, Docno: LA010189-0002

Prepare a query file or a dataframe for batch processing. This file should be in TREC format and should contain the following queries: "curbing population growth" (q435), "railway accidents" (q436), "inventions, scientific discoveries" (q439), "ship losses" (q448), "King Hussein, peace" (q450). The numbers in parentheses are query numbers. Show your query file or dataframe. [5 points]

```
queries_trec = pd.DataFrame([
    ["q435", "curbing population growth"],
    ["q436", "railway accidents"],
    ["q439", "inventions, scientific discoveries"],
    ["q448", "ship losses"],
    ["q450", "King Hussein, peace"]
], columns=["qid", "query"])
```

queries_trec

| | qid | query |
|---|------|------------------------------------|
| 0 | q435 | curbing population growth |
| 1 | q436 | railway accidents |
| 2 | q439 | inventions, scientific discoveries |
| 3 | q448 | ship losses |
| 4 | q450 | King Hussein, peace |

Perform retrieval using three different retrieval models/methods. Report the ranked results, up to 10, for each of these. [10 points]

```
index = pt.IndexFactory.of("c:/Users/Fortu/Downloads/Aut2024/Info 376/midterm_practice/la_times_in
BM = pt.BatchRetrieve(index, wmodel="BM25")
TF_IDF = pt.BatchRetrieve(index, wmodel="TF_IDF")
Hiem = pt.BatchRetrieve(index, wmodel="Hiemstra_LM")
```

C:\Users\Fortu\AppData\Local\Temp\ipykernel_9280\3012119167.py:2: DeprecationWarning: Call to deprecated class BatchRetrieve. (use pt.terrier.Retriever() instead) -- Deprecated since version 0.11.0.

```
BM = pt.BatchRetrieve(index, wmodel="BM25")
```

C:\Users\Fortu\AppData\Local\Temp\ipykernel_9280\3012119167.py:3: DeprecationWarning: Call to deprecated class BatchRetrieve. (use pt.terrier.Retriever() instead) -- Deprecated since version 0.11.0.

```
TF_IDF = pt.BatchRetrieve(index, wmodel="TF_IDF")
```

C:\Users\Fortu\AppData\Local\Temp\ipykernel_9280\3012119167.py:4: DeprecationWarning: Call to deprecated class BatchRetrieve. (use pt.terrier.Retriever() instead) -- Deprecated since version 0.11.0.

0.11.0.

```
Hiem = pt.BatchRetrieve(index, wmodel="Hiemstra_LM")
```

```
BM_values = BM.transform(queries_trec)
TF_IDF_values = TF_IDF.transform(queries_trec)
Hiem_values = Hiem.transform(queries_trec)
```

```
print(BM_values.head(10))
```

| | qid | docid | docno | rank | score | query |
|---|------|--------|---------------|------|-----------|---------------------------|
| 0 | q435 | 126171 | LA121589-0139 | 0 | 25.390996 | curbing population growth |
| 1 | q435 | 83083 | LA082189-0092 | 1 | 23.927200 | curbing population growth |
| 2 | q435 | 48392 | LA051590-0008 | 2 | 23.758792 | curbing population growth |
| 3 | q435 | 107074 | LA102689-0071 | 3 | 23.073475 | curbing population growth |
| 4 | q435 | 31179 | LA032990-0138 | 4 | 22.825540 | curbing population growth |
| 5 | q435 | 122235 | LA120589-0035 | 5 | 22.547941 | curbing population growth |
| 6 | q435 | 46844 | LA051189-0089 | 6 | 21.957725 | curbing population growth |
| 7 | q435 | 89431 | LA090890-0109 | 7 | 21.014325 | curbing population growth |
| 8 | q435 | 2032 | LA010790-0036 | 8 | 20.837134 | curbing population growth |
| 9 | q435 | 44580 | LA050490-0199 | 9 | 20.675308 | curbing population growth |

```
print(TF_IDF_values.head(10))
```

| | qid | docid | docno | rank | score | query |
|---|------|--------|---------------|------|-----------|---------------------------|
| 0 | q435 | 126171 | LA121589-0139 | 0 | 14.082744 | curbing population growth |
| 1 | q435 | 83083 | LA082189-0092 | 1 | 13.269034 | curbing population growth |
| 2 | q435 | 48392 | LA051590-0008 | 2 | 13.174410 | curbing population growth |
| 3 | q435 | 107074 | LA102689-0071 | 3 | 12.799793 | curbing population growth |
| 4 | q435 | 31179 | LA032990-0138 | 4 | 12.655082 | curbing population growth |
| 5 | q435 | 122235 | LA120589-0035 | 5 | 12.527306 | curbing population growth |
| 6 | q435 | 46844 | LA051189-0089 | 6 | 12.206764 | curbing population growth |
| 7 | q435 | 89431 | LA090890-0109 | 7 | 11.647640 | curbing population growth |
| 8 | q435 | 2032 | LA010790-0036 | 8 | 11.587420 | curbing population growth |
| 9 | q435 | 44580 | LA050490-0199 | 9 | 11.474252 | curbing population growth |

```
print(Hiem_values.head(10))
```

| | qid | docid | docno | rank | score | query |
|---|------|--------|---------------|------|-----------|---------------------------|
| 0 | q435 | 126171 | LA121589-0139 | 0 | 15.093337 | curbing population growth |
| 1 | q435 | 83083 | LA082189-0092 | 1 | 13.261315 | curbing population growth |
| 2 | q435 | 48392 | LA051590-0008 | 2 | 13.008375 | curbing population growth |
| 3 | q435 | 122235 | LA120589-0035 | 3 | 11.152714 | curbing population growth |
| 4 | q435 | 89431 | LA090890-0109 | 4 | 10.840118 | curbing population growth |
| 5 | q435 | 107074 | LA102689-0071 | 5 | 10.820350 | curbing population growth |
| 6 | q435 | 46844 | LA051189-0089 | 6 | 10.618389 | curbing population growth |
| 7 | q435 | 31179 | LA032990-0138 | 7 | 10.535393 | curbing population growth |
| 8 | q435 | 2032 | LA010790-0036 | 8 | 10.483613 | curbing population growth |
| 9 | q435 | 131639 | LA123189-0076 | 9 | 10.285353 | curbing population growth |

Use the provided [qrels document](#) to perform evaluation on these results. Report MAP, R-precision, and reciprocal rank values. Provide your brief thoughts on how well your retrieval models are doing. [15 points]

```
from pyterrier.measures import *

# Load the qrels file
qrels = pd.read_csv('C:/Users/Fortu/Downloads/Aut2024/Info 376/qrels.csv')

# Perform evaluation
results = pt.Experiment(
    [BM, TF_IDF, Hiem],
    queries_trec,
    qrels,
    ["map", "recip_rank", "Rprec"]
)

print(results)
```

| | name | map | recip_rank | Rprec |
|---|--------------------------|----------|------------|----------|
| 0 | TerrierRetr(BM25) | 0.021449 | 0.415009 | 0.044059 |
| 1 | TerrierRetr(TF_IDF) | 0.021637 | 0.414996 | 0.044742 |
| 2 | TerrierRetr(Hiemstra_LM) | 0.023595 | 0.425238 | 0.042602 |

Evaluation of Retrieval Models

The evaluation of the three retrieval models (BM25, TF_IDF, and Hiemstra_LM) using the metrics MAP, Reciprocal Rank, and R-precision provides insights into their performance:

- **BM25:**
 - **MAP:** 0.021449
 - **Reciprocal Rank:** 0.415009
 - **R-precision:** 0.044059
- **TF_IDF:**
 - **MAP:** 0.021637
 - **Reciprocal Rank:** 0.414996
 - **R-precision:** 0.044742
- **Hiemstra_LM:**
 - **MAP:** 0.023595
 - **Reciprocal Rank:** 0.425238
 - **R-precision:** 0.042602

Observations:

1. **MAP (Mean Average Precision):** Hiemstra_LM has the highest MAP, indicating it performs slightly better in terms of precision across all queries.
2. **Reciprocal Rank:** Hiemstra_LM also has the highest reciprocal rank, suggesting it is more effective at retrieving the top relevant document.

3. **R-precision:** TF_IDF has the highest R-precision, indicating it performs better in terms of precision at the rank cut-off point.

Overall, Hiemstra_LM appears to be the most effective model based on MAP and Reciprocal Rank, while TF_IDF shows its strength in R-precision (at least in this case). BM25 performs consistently but slightly lower than the other two models in these metrics. This shows that different models have varying strengths in retrieval performance, and the choice of model can impact the effectiveness of the retrieval system.